

# Statistical Machine Learning

## Introduction

# People

- ▶ Randy Lai (aka “Chu Shing Lai”)
- ▶ Office: Neville Hall 334
- ▶ Office Hour: MWF 11:00-12:00 (tentatively)

## Course Description

Focus on linear and nonlinear statistical models. Emphasis on concepts, methods, and data analysis; formal mathematics kept to minimum. Topics include resampling methods, regularization techniques in regression and modern classification, cluster analysis and dimension reduction techniques. Use professional level software.

# Logistics

- ▶ Course material (lectures slides, homeworks, etc.):  
<https://bb.courses.maine.edu>
- ▶ 5-6 assignments (work in groups!?)
- ▶ final project and presentation (kaggle competitions?)

# Text

- ▶ (Primary) *Introduction to Statistical Learning* by James, Witten, Hastie, and Tibshirani: download at <http://www-bcf.usc.edu/~gareth/ISL>
- ▶ *Computer Age Statistical Inference* by Efron and Hastie: download at <https://web.stanford.edu/~hastie/CASI/>
- ▶ *Elements of Statistical Learning* by Hastie, Tibshirani, and Friedman: download at <http://www-stat.stanford.edu/ElemStatLearn>

# Why you are here?

- ▶ Because you love the subject, because it's required, because you eventually want to make \$
- ▶ No matter the reason, (I wish) everyone can get something out of the course

# Outline

- ▶ Introduction
- ▶ Statistical Learning
- ▶ Linear Regression
- ▶ Classification
- ▶ Resampling Methods
- ▶ Linear Model Selection and Regularization
- ▶ Moving Beyond Linearity
- ▶ Tree-Based Methods
- ▶ Support Vector Machines
- ▶ Unsupervised Learning
- ▶ Neural Networks
- ▶ Anything else?

# A brief history of Statistical learning

- ▶ At the beginning of the nineteenth century, Legendre and Gauss published papers on the method *method of least square*
- ▶ Fisher proposed *linear discriminant analysis* in 1936
- ▶ In the 1940s, various authors introduced *logistic regression*
- ▶ In the early 1970s, Nelder and Wedderburn proposed generalized linear models
- ▶ In the late 70s/early 80s, Wahba published papers on *spline models*
- ▶ In mid 1980s Breiman, Friedman, Olshen and Stone introduced *classification and regression trees*



## A brief history of Statistical learning (cont')

- ▶ Hastie and Tibshirani coined the term *generalized additive models* in 1986
- ▶ Since then, inspired by the advent of *machine learning* and other disciplines, statistical learning has emerged as a new subfield in statistics, focused on *supervised* and *unsupervised* modeling and prediction

# Supervised Learning

Supervised learning: making predictions

- ▶ we have training data  $(X_1, Y_1), \dots, (X_n, Y_n)$  to learn a model in order to predict  $Y$  from  $X$
- ▶ outcome measurement  $Y$  (also called dependent variable, response, target).
- ▶ vector of  $p$  predictor measurements  $X$  (also called inputs, regressors, covariates, features, independent variables)
- ▶ In the regression problem,  $Y$  is quantitative
- ▶ In the classification problem,  $Y$  takes values in a finite, unordered set

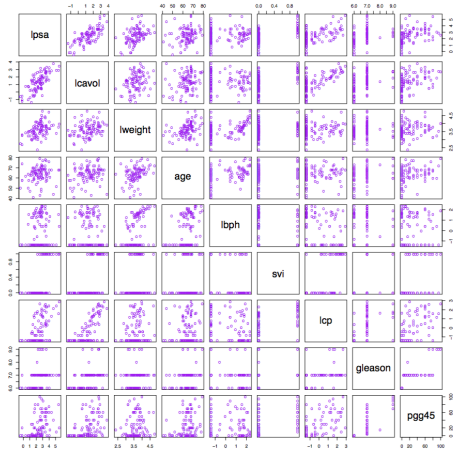
# Unsupervised learning

Unsupervised learning: discovering structure

- ▶ given measurements  $X_1, \dots, X_n$ , learn some underlying group structure based on similarity
- ▶ no outcome variable, just a set of predictors (features) measured on a set of samples
- ▶ objective is more fuzzy
- ▶ difficult to know how well you are doing.
- ▶ different from supervised learning, but can be useful as a pre-processing step for supervised learning


# Example 1

Identify the risk factors for prostate cancer.



## Example 2


Handwritten Digit Recognition: The data from this example come from the handwritten ZIP codes on envelopes from U.S. postal mail. Each image is a segment from a five digit ZIP code, isolating a single digit.



## Example 3




Classify a tissue sample into one of several cancer classes, based on a gene expression profile.

- ▶ which samples are most similar to each other?
- ▶ which genes are most similar to each other?
- ▶ do certain genes show very high (or low) expression for certain cancer samples?



## Example 4

Establish the relationship between salary and demographic variables in population survey data.



# Statistical Learning vs Machine Learning

- ▶ Machine learning arose as a subfield of Artificial Intelligence.
- ▶ Statistical learning arose as a subfield of Statistics.
- ▶ There is much overlap – both fields focus on supervised and unsupervised problems:
- ▶ Machine learning has a greater emphasis on large scale applications and prediction accuracy.
- ▶ Statistical learning emphasizes models and their interpretability, and precision and uncertainty.
- ▶ But the distinction has become more and more blurred, and there is a great deal of “cross-fertilization”.
- ▶ Machine learning has the upper hand in Marketing!