

Regression demo

```
library(tidyverse)
```

```
## -- Attaching packages -----
```

```
## v ggplot2 3.0.0      v purrr  0.2.5
## v tibble  1.4.2      v dplyr  0.7.6
## v tidyr   0.8.1      v stringr 1.3.1
## v readr   1.1.1      v forcats 0.3.0
```

```
## -- Conflicts -----
```

```
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

```
library(gapminder)
```

```
gapminder
```

```
## # A tibble: 1,704 x 6
```

	country	continent	year	lifeExp	pop	gdpPercap
	<fct>	<fct>	<int>	<dbl>	<int>	<dbl>
## 1	Afghanistan	Asia	1952	28.8	8425333	779.
## 2	Afghanistan	Asia	1957	30.3	9240934	821.
## 3	Afghanistan	Asia	1962	32.0	10267083	853.
## 4	Afghanistan	Asia	1967	34.0	11537966	836.
## 5	Afghanistan	Asia	1972	36.1	13079460	740.
## 6	Afghanistan	Asia	1977	38.4	14880372	786.
## 7	Afghanistan	Asia	1982	39.9	12881816	978.
## 8	Afghanistan	Asia	1987	40.8	13867957	852.
## 9	Afghanistan	Asia	1992	41.7	16317921	649.
## 10	Afghanistan	Asia	1997	41.8	22227415	635.

```
## # ... with 1,694 more rows
```

Example 1

We want to predict lifeExp using year and gdpPercap.

```
fit <- lm(lifeExp ~ year + gdpPercap, data = gapminder)
fit
```

```
##
```

```
## Call:
```

```
## lm(formula = lifeExp ~ year + gdpPercap, data = gapminder)
```

```
##
```

```
## Coefficients:
```

	year	gdpPercap
## (Intercept)		
##	-4.184e+02	2.390e-01 6.697e-04

confidence intervals of coefficients

```
confint(fit)
```

```
##                2.5 %      97.5 %  
## (Intercept) -4.725914e+02 -3.642571e+02  
## year        2.115805e-01  2.663851e-01  
## gdpPercap   6.217371e-04  7.177274e-04
```

```
confint(fit, "year")
```

```
##                2.5 %      97.5 %  
## year 0.2115805 0.2663851
```

```
confint(fit, 2)
```

```
##                2.5 %      97.5 %  
## year 0.2115805 0.2663851
```

```
confint(fit, level = 0.9)
```

```
##                5 %      95 %  
## (Intercept) -4.638752e+02 -3.729734e+02  
## year        2.159899e-01  2.619756e-01  
## gdpPercap   6.294602e-04  7.100043e-04
```

```
summary(fit)
```

```
##  
## Call:  
## lm(formula = lifeExp ~ year + gdpPercap, data = gapminder)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -67.262  -6.954   1.219   7.759  19.553   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept) -4.184e+02  2.762e+01  -15.15  <2e-16 ***  
## year        2.390e-01  1.397e-02   17.11  <2e-16 ***  
## gdpPercap   6.697e-04  2.447e-05   27.37  <2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 9.694 on 1701 degrees of freedom  
## Multiple R-squared:  0.4375, Adjusted R-squared:  0.4368   
## F-statistic: 661.4 on 2 and 1701 DF,  p-value: < 2.2e-16
```

```
library(broom)  
glance(fit)
```

```
## # A tibble: 1 x 11
##   r.squared adj.r.squared sigma statistic    p.value    df logLik    AIC
## *   <dbl>         <dbl> <dbl>    <dbl>    <dbl> <int>  <dbl>  <dbl>
## 1    0.437         0.437  9.69    661. 3.13e-213     3 -6287. 12582.
## # ... with 3 more variables: BIC <dbl>, deviance <dbl>, df.residual <int>
```

Comparing models via adjusted R^2

```
fit0 <- lm(lifeExp ~ year, data = gapminder)
fit2 <- lm(lifeExp ~ year + pop, data = gapminder)
fit3 <- lm(lifeExp ~ year + gdpPercap + pop, data = gapminder)
```

```
bind_rows(
  glance(fit0),
  glance(fit),
  glance(fit2),
  glance(fit3)
)
```

```
## # A tibble: 4 x 11
##   r.squared adj.r.squared sigma statistic    p.value    df logLik    AIC
##   <dbl>         <dbl> <dbl>    <dbl>    <dbl> <int>  <dbl>  <dbl>
## 1    0.190         0.189 11.6    399. 7.55e- 80     2 -6598. 13202.
## 2    0.437         0.437  9.69    661. 3.13e-213     3 -6287. 12582.
## 3    0.191         0.190 11.6    200. 7.73e- 79     3 -6597. 13202.
## 4    0.440         0.439  9.67    446. 1.52e-213     4 -6283. 12576.
## # ... with 3 more variables: BIC <dbl>, deviance <dbl>, df.residual <int>
```

Assessing the Accuracy of the estimated regression model

```
library(modelr)
```

```
##
## Attaching package: 'modelr'
```

```
## The following object is masked from 'package:broom':
##
##   bootstrap
```

```
rp <- resample_partition(gapminder, c(train = 0.7, test = 0.3))
rp
```

```
## $train
## <resample [1,192 x 6]> 1, 3, 4, 5, 7, 8, 9, 12, 13, 14, ...
##
## $test
## <resample [512 x 6]> 2, 6, 10, 11, 16, 22, 23, 26, 39, 40, ...
```

```
training_set <- as.tibble(rp$train)
testing_set <- as.tibble(rp$test)
```

```
fit0 <- lm(lifeExp ~ year, data = training_set)
fit1 <- lm(lifeExp ~ year + gdpPercap, data = training_set)
fit2 <- lm(lifeExp ~ year + pop, data = training_set)
fit3 <- lm(lifeExp ~ year + gdpPercap + pop, data = training_set)
```

```
c(mse(fit0, testing_set),
  mse(fit1, testing_set),
  mse(fit2, testing_set),
  mse(fit3, testing_set))
```

```
## [1] 137.3380 104.3193 137.0711 103.7084
```

Interpreting regression coefficients

```
supermodel <- read_tsv("supermodel.dat")
```

```
## Parsed with column specification:
## cols(
##   salary = col_double(),
##   age = col_double(),
##   years = col_double(),
##   beauty = col_double()
## )
```

```
supermodel
```


```
## # A tibble: 231 x 4
##   salary    age years beauty
##   <dbl> <dbl> <dbl> <dbl>
## 1  0.370   16.7  3.15  78.3
## 2 53.7    20.3  5.51  68.6
## 3  1.46    18.2  5.33  75.0
## 4 0.0243   15.4  3.84  65.1
## 5 95.3     24.2  8.53  71.8
## 6 14.6     18.3  4.39  78.1
## 7  8.67    17.7  4.40  72.1
## 8  2.65    17.5  4.11  75.3
## 9  7.55    17.1  3.52  72.0
## 10 1.20     20.1  6.83  71.9
## # ... with 221 more rows
```

```
(fit <- lm(salary ~ age + years + beauty, data = supermodel))
```

```
##
## Call:
```

```
## lm(formula = salary ~ age + years + beauty, data = supermodel)
##
## Coefficients:
## (Intercept)      age      years      beauty
##   -60.8897    6.2344   -5.5612   -0.1964
```

```
ggplot(supermodel) + geom_point(aes(x = age, y = years))
```




Correlations among predictors

```
x1 <- rnorm(100)
x2 <- 2 * x1 + rnorm(100, sd = 0.1)
y <- 3 + 2 * x1 + rnorm(100, sd = 0.5)
(example <- tibble(y = y, x1 = x1, x2 = x2))
```


```
## # A tibble: 100 x 3
##       y      x1      x2
##   <dbl> <dbl> <dbl>
## 1 5.64  0.905  1.90
## 2 3.02  0.0620 0.176
## 3 1.32 -0.812 -1.66
## 4 3.91  0.624  1.17
## 5 2.13 -0.389 -0.885
```

```
## 6 0.758 -0.889 -1.87
## 7 1.64 -1.05 -2.01
## 8 2.05 -0.308 -0.707
## 9 2.22 -0.530 -0.874
## 10 1.99 -0.419 -0.914
## # ... with 90 more rows
```

```
ggplot(example) + geom_point(aes(x1, y))
```



```
ggplot(example) + geom_point(aes(x1, x2))
```



```
(fit <- lm(y ~ x1 + x2, data = example))
```

```
##
## Call:
## lm(formula = y ~ x1 + x2, data = example)
##
## Coefficients:
## (Intercept)      x1      x2
##      3.0164    0.5190    0.7541
```

Non linear transformation of the predictors

```
rp <- resample_partition(gapminder, c(train = 0.7, test = 0.3))
rp
```

```
## $train
## <resample [1,192 x 6]> 1, 2, 3, 6, 7, 8, 10, 11, 15, 18, ...
##
## $test
## <resample [512 x 6]> 4, 5, 9, 12, 13, 14, 16, 17, 22, 24, ...
```

```

training_set <- as.tibble(rp$train)
testing_set <- as.tibble(rp$test)
fit_linear <- lm(lifeExp ~ gdpPercap, data = training_set)
fit_quad <- lm(lifeExp ~ gdpPercap + I(gdpPercap^2), data = training_set)

```

Which one is better?

```
c(mse(fit_linear, testing_set), mse(fit_quad, testing_set))
```

```
## [1] 100.33030 80.03267
```

How about adding a cubic term? Adding $I(\text{gdpPercap}^3)$ term? No. Use `poly`!

```

fit_cubic <- lm(lifeExp ~ poly(gdpPercap, 3), data = training_set)
mse(fit_cubic, testing_set) # even better


```

```
## [1] 67.38276
```

```

ggplot(gapminder, aes(x = gdpPercap, y = lifeExp)) + geom_point() +
  geom_smooth(method = "lm", formula = y ~ poly(x, 10))

```



Let's be crazy


```
fit_poly <- lm(lifeExp ~ poly(gdpPercap, 8), data = training_set)
mse(fit_poly, testing_set) # becoming worse
```

```
## [1] 54.29174
```


Of course, we are not limited to polynomials

```
fit_log <- lm(lifeExp ~ log(gdpPercap), data = training_set)
mse(fit_log, testing_set) # wow, even better
```

```
## [1] 56.3479
```

In fact, look at this

```
ggplot(gapminder, aes(x = gdpPercap, y = lifeExp)) + geom_point() +
  geom_smooth(method = "lm", formula = y ~ log(x))
```



Qualitative Predictor

```
gapminder %>% distinct(country)
```

```
## # A tibble: 142 x 1
##   country
##   <fct>
## 1 Afghanistan
## 2 Albania
## 3 Algeria
## 4 Angola
## 5 Argentina
## 6 Australia
## 7 Austria
## 8 Bahrain
## 9 Bangladesh
## 10 Belgium
## # ... with 132 more rows
```

Let's consider a smaller

```
little_gapminder <- gapminder %>% filter(str_detect(country, "^H")) %>%
  mutate(country = fct_drop(country))
little_gapminder %>% distinct(country)
```

```
## # A tibble: 4 x 1
##   country
##   <fct>
## 1 Haiti
## 2 Honduras
## 3 Hong Kong, China
## 4 Hungary
```

Now, we want to use `country` to predict `lifeExp`.

```
(fit <- lm(lifeExp ~ country, data = little_gapminder))
```

```
##
## Call:
## lm(formula = lifeExp ~ country, data = little_gapminder)
##
## Coefficients:
##           (Intercept)          countryHonduras  countryHong Kong, China
##                50.165                7.756                23.328
##          countryHungary
##                19.228
```

Under the hood.

```
contrasts(little_gapminder$country)
```

```
##           Honduras Hong Kong, China Hungary
## Haiti           0           0           0
## Honduras         1           0           0
## Hong Kong, China  0           1           0
## Hungary           0           0           1
```

A common mistake

```
(example2 <- tibble(x = sample(1:3, 10, replace = TRUE), y = rnorm(10)))
```

```
## # A tibble: 10 x 2
##       x         y
##   <int>   <dbl>
## 1     2  1.42
## 2     1  0.0429
## 3     1 -0.456
## 4     3 -0.518
## 5     2 -0.840
## 6     3  0.860
## 7     3 -1.31
## 8     3 -0.239
## 9     2  0.420
## 10    2  0.991
```

```
lm(y ~ x, data = example2) # wrong
```

```
##
## Call:
## lm(formula = y ~ x, data = example2)
##
## Coefficients:
## (Intercept)              x
##      0.3785      -0.1551
```

```
example2_corrected <- example2 %>% mutate(x = recode_factor(x, `1` = "blue", `2` = "green", `3` = "red"))
example2_corrected
```

```
## # A tibble: 10 x 2
##       x         y
##   <fct>   <dbl>
## 1 green  1.42
## 2 blue   0.0429
## 3 blue  -0.456
## 4 red   -0.518
## 5 green -0.840
## 6 red    0.860
## 7 red   -1.31
## 8 red   -0.239
## 9 green  0.420
## 10 green 0.991
```

```
lm(y ~ x, data = example2_corrected) # correct
```

```
##
## Call:
## lm(formula = y ~ x, data = example2_corrected)
```

```
##
## Coefficients:
## (Intercept)      xgreen      xred
##      -0.20632      0.70426     -0.09535
```

```
contrasts(example2_corrected$x)
```

```
##      green red
## blue      0  0
## green      1  0
## red        0  1
```

Prediction

```
fit <- lm(lifeExp ~ year, data = gapminder)
new_data <- tibble(year = 100)
```

```
# the classic way
predict(fit, new_data)
```

```
##      1
## -553.0618
```

```
# in tidyverse style
library(modelr)
new_data %>% add_predictions(fit)
```

```
## # A tibble: 1 x 2
##   year pred
##   <dbl> <dbl>
## 1   100 -553.
```